Modeling Soccer Match Probabilities and Beating the Bookmaker: Statistical Methods, Machine Learning, and Pricing Mechanics

Pablo Sm · June 21, 2026

Executive Summary

The science of forecasting soccer matches has matured along two parallel tracks that this report synthesizes: a **classical statistical lineage** running from Maher's 1982 independent Poisson model through Dixon-Coles (1997), bivariate Poisson (Karlis-Ntzoufras 2003), and dynamic latent-strength frameworks (Rue-Salvesen 2000), and a **modern machine-learning track** built on gradient-boosted trees, neural networks, and ensemble systems. Both tracks ultimately serve the same commercial purpose: estimating the probability matrix of match scorelines so that win/draw/loss and derivative-market prices can be generated, compared against bookmaker lines, and exploited where mispricing exists [1][2][3]. The foundational insight that unifies the field is that goals are well-approximated as Poisson-distributed counts driven by team attack strength, opponent defense strength, and home advantage — but that the naive independence assumption systematically underestimates low-scoring draws, a flaw every serious model since 1997 has been engineered to correct [1][4].

The single most important practical finding across the research is that **calibration matters more than accuracy for profitability**. Studies using NBA data found that models selected for calibration achieved an average ROI of +34.69%, while models selected for raw accuracy produced a negative ROI of -35.17% [5][6]. This is decisive for anyone building a betting model: predicting the winner correctly is worthless if the predicted probabilities are miscalibrated, because every value calculation and every Kelly stake depends on the probability number, not the classification label [7][8]. The bookmaker's closing line is the most efficient consensus estimate available, and Closing Line Value (CLV) is the field's preferred process-based benchmark for distinguishing genuine edge from luck [9][10][11]. Markets are highly — but not perfectly — efficient: the favorite-longshot bias persists, top-tier leagues are sharper than lower tiers, and arbitrage appears in well under 0.5% of matches, meaning realistic edges come from statistical mispricing rather than pure arbitrage [12][13][14].

For decision-makers, the path to a defensible custom model is now well-charted but unforgiving. It requires walk-forward backtesting that rigorously avoids data leakage, devigging of bookmaker odds to recover true implied probabilities, post-hoc calibration (Platt scaling, isotonic regression, or Beta calibration), evaluation by Ranked Probability Score and log-loss rather than accuracy, and disciplined fractional-Kelly staking to survive variance [15][16][17][7]. Even so, the research is sober about the odds of success: bookmakers continuously improve forecast precision under competitive pressure, the overround is a structural barrier, and benchmarks like "KellyBench" show that even frontier AI models often fail to turn a profit over a full season [18][19].

> **Bottom line for decision-makers:** The mathematics of soccer modeling is mature and openly available; the durable edge lies in disciplined execution — leakage-free data pipelines, probability calibration, devigging, CLV measurement, and fractional-Kelly risk control. Treat the closing line as the benchmark to beat, prioritize calibration over accuracy, and assume the market is your toughest competitor, not your customer.

Key Findings

1. **Calibration, not accuracy, drives betting returns.** Models selected for calibration delivered +34.69% ROI versus -35.17% for models selected for accuracy in head-to-head NBA comparisons [5][6]. **So what:** Build, tune, and select your model on proper scoring rules (Brier, log-loss, RPS) and reliability diagrams — never on classification accuracy — because every value and stake decision is a function of the probability number itself [7][8].

2. **The Dixon-Coles correction remains the gold-standard fix for the Poisson model's central flaw.** The independent Poisson model systematically underestimates low-scoring draws (0-0, 1-1), which Dixon-Coles repairs with a correlation parameter (rho) and a time-decay weighting of recent matches [1][4][20]. **So what:** Any goals-based model you build should start from Dixon-Coles as a baseline benchmark; if a more complex model cannot beat it on RPS, the complexity is not earning its keep [4].

3. **Bookmaker odds are the strongest single predictor available, and ignoring them is a measurable handicap.** Models that fail to account for market prices operate at a Kullback-Leibler disadvantage relative to the market [21][22][23]. **So what:** Incorporate market-implied probabilities directly — via Bayesian combination with your historical attack-defense estimates — rather than trying to beat the market from historical data alone [22].

4. **Devigging is a non-trivial modeling step that materially affects measured edge.** Bookmaker odds are polluted by margin, favorite-longshot bias, and risk policy; methods range from naive normalization to the Shin method and domain-driven multi-market integration [16]. **So what:** Choose a devig method deliberately and test it — naive normalization can over- or under-state true probabilities because margins are not applied uniformly across favorites and longshots [24][25].

5. **The favorite-longshot bias persists in fixed-odds soccer markets but is league- and tier-dependent.** Longshots are systematically over-bet and favorites under-bet; top-tier leagues (e.g., markets priced by Pinnacle) are more efficient than lower tiers [13][26][14]. **So what:** Concentrate exploratory edge-hunting in less liquid, lower-tier, or derivative markets where pricing is softer, rather than in the most heavily traded top-league 1X2 markets [13][27].

6. **Closing Line Value (CLV) is the most reliable indicator of long-term skill.** Because single results are dominated by variance, beating the closing price consistently is a better signal of edge than short-run profit and loss [9][11][28]. **So what:** Instrument your pipeline to record the closing line for every bet and track CLV as your primary KPI — it tells you whether your process works long before your bankroll does [28].

7. **Data leakage is the dominant failure mode in soccer modeling.** Using half-time results to predict full-time outcomes, or including a game's own result in a team's running season totals, are classic leaks that inflate backtest performance [17][29]. **So what:** Adopt strict walk-forward training where features use only information available before kickoff, and audit every cumulative feature for self-inclusion [27][29].

8. **Ensembles and gradient-boosted trees consistently outperform single models on tabular soccer data.** XGBoost and CatBoost paired with rating features (Elo, pi-ratings) are repeatedly cited as top performers, and voting/stacking ensembles beat individual models [20][30][31][32]. **So what:** Prefer a calibrated ensemble of a goals-based statistical model and a gradient-boosted tree model over any single algorithm, then calibrate the ensemble output [32][31].

9. **Fractional Kelly is the practical standard for staking.** Full Kelly assumes perfect probability knowledge and produces extreme volatility; quarter- or half-Kelly reduces variance and the impact of estimation error [15][27][33]. **So what:** Never stake full Kelly with model-estimated (rather than true) probabilities; the volatility and risk of ruin are not survivable for human bankrolls [34][35].

10. **Even sophisticated models rarely beat the market over a full season.** The "KellyBench" benchmark showed frontier AI models often fail to achieve profitability across a season, reflecting non-stationarity and the overround barrier [19]. **So what:** Set realistic expectations — thin, intermittent edges in specific sub-markets are the realistic prize, not a reliable income stream [19].

Detailed Analysis

The Statistical Lineage: From Poisson Counts to Correlated Scorelines

The entire edifice of soccer modeling rests on a single observation: goals are rare, discrete events that arrive at a roughly constant rate, making them well-suited to the Poisson distribution [36][1]. The seminal application is **M. J. Maher's 1982 paper "Modelling association football scores"**, published in Statistica Neerlandica, which proposed that the goals of the two teams could be modeled as independent Poisson random variables [37][38][2][39]. In this framework the home team's goals X follow Poisson(lambda_h) and the away team's goals Y follow Poisson(lambda_a), with the rate parameters reflecting the attacking and defensive strengths of the competing teams, estimated by maximum likelihood [1][37][38][40].

This produces a complete scoreline probability matrix. From it, one derives the probability of any match outcome — home win, draw, away win — by summing the appropriate cells, and of any derivative market (over/under goals, both-teams-to-score, correct score) by aggregating differently [38][40]. The parameters are typically fit using a **Generalized Linear Model (GLM) with a log-link function**, where the log of the expected goals is a linear combination of an attack parameter for the scoring team, a defense parameter for the conceding team, and a home-advantage term [1][20].

Maher himself acknowledged the model's central weakness: the independent Poisson assumption does not perfectly capture match data, particularly the observed frequency of draws [38]. This single shortcoming launched four decades of refinement.

The Independence Problem and the Dixon-Coles Repair

The most influential extension is the **Dixon-Coles model**, introduced in 1997 by Mark J. Dixon and Stuart G. Coles in the Journal of the Royal Statistical Society (Series C) [1]. Their motivation was precisely the empirical failure of independent Poisson: assuming home and away goals are independent leads to systematic bias, especially in low-scoring games where the model underestimates the frequency of draws such as 0-0 and 1-1 [1][4].

Dixon-Coles makes two signature modifications:

- **A correlation parameter (rho)** — an interaction term that adjusts the joint probability mass function for the low-scoring scorelines (0-0, 1-0, 0-1, 1-1), increasing their likelihood to better match observed data [1][4][20].
- **Time-weighting** via an exponential decay function, often written exp(-xi * t), which assigns greater weight to recent matches so the estimated team strengths reflect current form rather than stale results [20][1][41].

The model retains the attack parameter (alpha), defense parameter (beta), and home-advantage parameter (gamma) of the Poisson framework, computing expected home goals (lambda_h) and away goals (lambda_a) from them, and is estimated by maximum likelihood [20][1]. Because it models the full scoreline matrix, it serves not only the 1X2 market but also Both-Teams-To-Score, Asian handicaps, and goal totals [42][43]. It is implemented in open-source libraries including the R package `goalmodel` and the Python package `penaltyblog` [42][20][43].

The research describes Dixon-Coles as the "gold standard" in academic sports analytics, while noting clear limitations: it does not inherently account for tactical systems or live match dynamics [4]. This has driven extensions toward bivariate Poisson models and machine-learning-integrated approaches [44][1].

Modeling Correlation Directly: Bivariate Poisson and Skellam

Where Dixon-Coles patches the independence assumption with a local correction term, **Dimitris Karlis and Ioannis Ntzoufras** addressed it structurally in their 2003 paper "Analysis of Sports Data Using Bivariate Poisson Models" [3][45][46]. They model the joint distribution of home and away goals as a **bivariate Poisson distribution**, which explicitly incorporates the dependence between the scoring performances of opposing teams [3][47][45].

Key features of the bivariate Poisson approach:

- **Correlation handling** is built into the distribution itself rather than added as a correction [3].
- **Overdispersion** — where the sample variance exceeds the sample mean, a common phenomenon in football goal counts — is accommodated [47].

- **Diagonal inflation** extends the basic model to fit the elevated frequency of draws, allowing a better fit for tied results [47][46].
- **Bayesian estimation** is used, allowing covariates such as team strength, historical performance, and home advantage to enter via prior distributions [47][48].

Notably, Karlis and Ntzoufras later shifted toward modeling the **difference** in goals (the margin of victory) using the **Skellam distribution** [47][48]. This offers two distinct advantages: it eliminates the need to explicitly model the correlation between teams, and it avoids the assumption that individual team scores must be marginally Poisson distributed [47][48]. The Skellam approach remains in active use; modern forecasting frameworks employ it to model the difference between two Poisson scoring processes as a robust basis for win-draw-loss probabilities [15].

Dynamic Team Strength: The Rue-Salvesen Contribution

The Poisson and bivariate Poisson models treat team strength as fixed within an estimation window. **Håvard Rue and Øyvind Salvesen's 2000 paper "Prediction and retrospective analysis of soccer matches in a league"** broke from this by treating team strengths as **latent variables that evolve over time**, using Markov Chains to capture fluctuations in "team form" [49]. This made quality a dynamic variable rather than a static parameter, capturing momentum across a season [49].

Critically, Rue and Salvesen compared their model-based predictions against the fixed-odds offered by commercial bookmakers (such as Intertops) to assess market efficiency and predictive accuracy — an early instance of the model-versus-market discipline that defines value betting today [44][30][1]. The model is now a standard benchmark against which newer machine-learning techniques like Random Forests, XGBoost, and neural networks are contrasted [49], and its handling of latent, time-varying strength remains a reference point for time-series sports analytics [49][44].

Bayesian and Distributional Extensions

Bayesian methodology has become increasingly prevalent because it provides a flexible framework for parameter estimation and uncertainty quantification [36][50]. The research identifies several strands:

- **Hierarchical models** estimate team abilities as latent variables and can incorporate prior information such as bookmaker odds to improve predictive accuracy [36][51].
- **Zero-modified / zero-inflated Poisson** distributions handle count data where the frequency of zero-goal outcomes deviates from standard Poisson expectations, adding flexibility [36].
- **Negative binomial and Weibull** approaches address overdispersion and model inter-arrival times of goals to better capture the "nature of the game" [36].
- **MCMC simulation** is used within Bayesian frameworks to derive outcome probabilities and tournament-level events such as qualification or relegation [50].

The most consequential modern refinement is the incorporation of **Expected Goals (xG)** into the underlying scoring rates (lambda), moving beyond final scores to reflect the quality of chances created [52][36]. xG is treated as a more stable performance metric than raw goal counts because it filters out finishing luck [15][27].

The table below summarizes the classical lineage.

Model	Year / Authors	Core innovation	Key weakness addressed	Estimation
Independent Poisson	1982, Maher	Goals as independent Poisson, attack/defense/ home params	First rigorous probabilistic framework	MLE [37][40]
Dixon-Coles	1997, Dixon & Coles	rho correlation term + time decay	Underestimation of low-scoring draws; stale form	MLE [1][20]
Bivariate Poisson	2003, Karlis & Ntzoufras	Joint distribution with correlation + diagonal inflation	Independence assumption; overdispersion; draws	Bayesian [3] [47]
Skellam (goal difference)	Karlis & Ntzoufras (later)	Models margin directly	Removes need to model correlation / marginal Poisson	Bayesian [47][48]
Rue-Salvesen	2000	Latent, time-evolving team strength via Markov chains	Static strength assumption; form/ momentum	Markov chain [49]
Bayesian hierarchical / zero-modified	Recent	Latent abilities, priors from odds, zero-inflation	Uncertainty quantification; excess zeros	MCMC [36] [50]

> **Bottom line for this theme:** Forty years of statistical refinement all circle the same problem — the independent Poisson model misprices draws and ignores form. Dixon-Coles remains the pragmatic baseline; bivariate Poisson and Skellam handle correlation more rigorously; Rue-Salvesen and Bayesian hierarchical models add dynamics. Start simple, benchmark everything against Dixon-Coles, and only add complexity that demonstrably lowers RPS [4].

The Machine Learning Track: Trees, Networks, and Ensembles

Soccer outcome prediction has evolved from traditional statistical models toward sophisticated machine-learning and deep-learning frameworks [49][30]. The research is consistent on one hierarchy of performance for **tabular** match data: gradient-boosted trees lead, neural networks are rising but situational, and ensembles beat individuals.

Tree-Based Methods Dominate Tabular Data

Gradient-boosted tree models — particularly **XGBoost** and **CatBoost** — are consistently cited as top performers, frequently paired with rating-based features such as pi-ratings and Elo to achieve superior accuracy [20][30][1]. These models iteratively improve accuracy by focusing on the errors of previous weak learners [4]. **Random Forest** is valued for robustness, its ability to handle complex feature interactions, embedded feature importance, and variance reduction through aggregation of many decision trees [42][30][4]. XGBoost in particular is a standard benchmark that effectively handles tabular match statistics and historical performance metrics, often outperforming simpler models, and has even been applied to detect tactical events like "line breaks" from player position and velocity data [53][49][20][41].

The research repeatedly stresses that **performance varies with feature set and tournament context** — Random Forest and Gradient Boosting are both effective, but which wins depends on the specific data and competition [44][4].

Neural Networks and Deep Learning: Powerful but Situational

Deep learning is increasingly used to model the temporal dynamics and player interactions that tabular models miss [30][43]:

- **Feedforward Neural Networks (FNNs)** have shown success on specific league datasets [1].
- **Recurrent Neural Networks (RNNs)** and **LSTMs** capture temporal dependencies such as team form and historical performance trends [30][44][4].
- **Convolutional Neural Networks (CNNs)** capture spatial/temporal patterns [53][44][20].
- **Transformers** with self-attention model long-range dependencies in match data, and **TimesNet-inspired** and **axial transformer** modules process complex multi-dimensional tracking data for in-game outcome forecasting [53][44][20][1][43].

These architectures have demonstrated competitive performance against traditional statistical and boosted-tree models [43], but the research does not claim they dominate on standard tabular features — their advantage appears greatest with high-granularity tracking and time-series data.

Ensembles: The Consistent Winner

The clearest practical recommendation from the ML literature is to combine models. **Stacking and voting ensembles** — aggregating predictions from logistic regression, random forests, XGBoost, AdaBoost, k-nearest neighbors, and others — consistently outperform single-model baselines [53][1][31][54]. Ensemble methods mitigate overfitting and bias inherent in single-model approaches by leveraging diverse architectures, reducing variance, and improving generalizability across leagues [31][1][32]. Both majority (hard) voting and soft (probability-averaging) voting are standard strategies [42][30].

Player- and Event-Level Feature Engineering

The research is emphatic that **feature quality often matters more than algorithm complexity** [30][15][55][27]. Feature categories cited include:

- **Historical results**: match results, goal differences, team rankings [42][30][4].
- **Player-level data**: goals, assists, passing accuracy, defensive actions, FIFA ratings — providing deeper insight into team synergy than aggregate stats [42][30][53].
- **Recency features**: weighted performance in recent matches [53][44][20][1].
- **Contextual data**: weather, venue, half-time status [53][44][20][1].
- **Network metrics**: passing networks where players are nodes and passes are weighted edges, with topological metrics (degree, closeness, betweenness, eigenvector centrality) capturing team strategy and synergy more effectively than box-score statistics alone [49][1].

Supporting techniques include **dimensionality reduction** via PCA to manage high-dimensional data and mitigate overfitting [42][20], and **data augmentation** such as SVM-SMOTE and Near-Miss to address class imbalance (e.g., the relative rarity of certain outcomes) [1].

The Expected Goals Revolution

xG modeling deserves separate treatment because it has reshaped feature engineering. xG represents the probability that a specific shot results in a goal, a more objective measure than total goals or raw shot counts [53][30]. Early models used only shot location and angle and were criticized for ignoring context such as defender positioning [53][20]. Contemporary models incorporate:

- **Defensive context**: number of nearby opponents, defensive pressure, goalkeeper positioning relative to ball and goal [20].
- **Spatiotemporal build-up**: the dynamics of the ~10-second sequence before a shot, using high-resolution event data (e.g., StatsBomb) and tracking data [53][30].
- **Player and position adjustment**: "player-adjusted" and "position-adjusted" xG segment data into Forward/Midfield/Defense groups and quantify individual efficiency — demonstrating that elite players (the research cites Messi as an example) consistently outperform the average xG model, proving individual quality is a critical and often underestimated variable [53][42][49].

Methodologically, xG models span **logistic regression, Random Forest, Gradient Boosting, and Multilayer Perceptrons** [53][42][20][30], with newer work using **Bayesian mixed-effects models** to add transparency and random effects for inter-player and inter-team variability [41], **Graph Neural Networks** to model players and ball as nodes capturing coordinated behaviors [1][43], and **offline Reinforcement Learning** to identify "optimal" actions and compare actual decisions against an optimal policy [30][43]. New frameworks even model **shot creation (Expected Threat / xS) jointly with conversion (xG)** so that chance builders are credited even when not the final finisher [44].

A persistent theme is the **interpretability-versus-performance trade-off**: black-box deep models yield high accuracy but low transparency, prompting use of explainable-AI tools like **SHAP** to isolate feature impact, and even “wordalization” via Large Language Models to translate regression coefficients into natural-language narratives for coaching staff [20][41][43].

Approach	Best use case	Strengths	Weaknesses
Random Forest	Baseline tabular classification	Robust, low overfitting, feature importance	Can plateau vs boosting [42][30][4]
XGBoost / CatBoost	Tabular match + rating features	Top accuracy, efficient, handles interactions	Needs calibration; context-dependent [20][30][43]
FNN / RNN / LSTM	Temporal form, sequences	Captures time dependencies	Data-hungry, harder to calibrate [30][44]
Transformers / TimesNet	Tracking & in-game time series	Long-range dependencies, in-play forecasting	Requires high-granularity data [44][1][43]
GNNs	Passing/relational structure	Captures synergy, coordination	Complex, data-intensive [1][43]
Bayesian mixed-effects xG	Player/position-adjusted xG	Transparent, quantifies variability	Less raw accuracy than black boxes [41]
Voting / stacking ensemble	Production forecasting	Best accuracy & robustness	Added complexity; still needs calibration [31][32]

> **Bottom line for this theme:** For tabular soccer data, a calibrated ensemble anchored by gradient-boosted trees and enriched with rating and xG features is the pragmatic frontier. Deep learning earns its place only when you have tracking or event-level time-series data. Spend your effort on feature engineering and calibration, not on chasing the most exotic architecture [15][31].

Bookmaker Pricing Mechanics: Overround, Vig, and the Margin Machine

Understanding how books price is prerequisite to beating them. The core fact: **bookmakers do not offer fair odds.** They quote odds whose implied probabilities sum to more than 100%, and this excess is variously called the booksum, vig, juice, or overround [56][57].

The Arithmetic of the Margin

The mechanics are precise and worth stating in plain terms:

- **Implied probability** = 1 divided by the decimal odds [58][59][60].
- **Overround** = (sum of all implied probabilities) minus 100% [58][59].
- **Vigorish** as a percentage of turnover = (overround divided by sum of implied probabilities) times 100 [59].

This margin is the bookmaker's built-in profit buffer, ensuring the total payout to winners is less than the total wagered [24][61]. The research also distinguishes three related but separate concepts that practitioners conflate at their peril [62]:

- **House edge (theoretical margin)**: the expected profit built into the odds.
- **Overround**: the excess of summed implied probabilities over 100% — the operator's cushion.
- **Hold (realized profit)**: the actual percentage of stakes retained after settlement.

A crucial structural point for multi-leg bettors: in **parlays and multiples, the overround compounds multiplicatively rather than additively**, sharply increasing the bookmaker's margin on these products [58][59]. **So what:** Avoid parlays as a value strategy — the margin stacks against you with each leg [58].

Odds Generation: From Compiler to Closing Line

Odds compilers (traders) set opening lines using a blend of quantitative modeling and qualitative judgment [59][63]. Sophisticated operators use simulations, Elo-style ratings, and Poisson distribution models to estimate true probabilities from team strength, form, injuries, and historical data [59][63]. The **opening line** is a starting point designed to attract liquidity; as the event approaches, the price moves toward the **closing line**, which is considered the most informationally efficient price because it has incorporated all available data and betting volume [59][63].

The book's job is not pure prediction but **risk management** [58][63]. If betting volume is lopsided, the operator adjusts odds to encourage action on the opposite side and minimize exposure — "balancing the book" [59][63]. Odds also move in response to **sharp money** (professional bettors and syndicates) whose wagers signal high-value information, while **public money** follows narratives, enabling "shading" where books adjust prices to exploit predictable public bias [59][63].

The Favorite-Longshot Bias as a Pricing Feature

The most documented systematic deviation is the **favorite-longshot bias (FLB)**: bettors overvalue longshots and undervalue favorites, so the implied win probabilities for longshots are too high and for favorites too low [18][26][14][64]. This is not necessarily bettor irrationality. The research offers structural explanations:

- **Bookmaker risk aversion**: books face higher profit variance on longshot bets, so they shade odds to mitigate risk, producing an FLB-like pattern even in competitive fixed-odds markets [64].
- **Imperfect competition and heterogeneous beliefs**: when bettors hold diverse beliefs, demand for longshots is less price-sensitive, letting books set odds that deviate from fair value [65].
- **Insider models (Shin)**: Shin's 1991/1992 models attribute FLB to the presence of insiders or disagreement among bettors; modern work suggests insiders' impact on the favorite-to-longshot odds ratio is often minimal, and markets may collapse if the insider fraction grows too large [66].

The FLB has a direct consequence for margin estimation: because books apply **higher margins to lower-probability (longshot) outcomes**, standard overround calculations that assume uniform margins **overestimate the expected payout** to bettors — meaning real loss rates are worse than the quoted overround implies [24][61][67][68][25].

Pricing concept	Definition	Practical implication
Implied probability	1 / decimal odds	Raw, margin-inflated probability [58]
Overround	Sum of implied probs − 100%	Operator cushion; varies by market [59]
Vig	Overround / sum of implied probs × 100	Margin as % of turnover [59]
Hold	Actual % of stakes retained	Realized, not theoretical, profit [62]
FLB margin skew	Higher margin on longshots	True loss rate > quoted overround [24][25]
Parlay compounding	Margins multiply across legs	Multiples carry the worst value [58]

> **Bottom line for this theme:** The overround is a structural barrier, not a constant — it is heavier on longshots and compounds in parlays. To recover usable probabilities you must devig deliberately, and you must assume your true loss rate is worse than the headline margin suggests [24][25].

Market Efficiency, Line Movement, and Where Edges Hide

The research presents a nuanced and sometimes contradictory picture of how efficient soccer betting markets really are — and the disagreements themselves are instructive.

How Efficient Are These Markets?

The consensus is that markets are **highly but not perfectly efficient**. Betting markets serve as "real-world laboratories" for the Efficient Market Hypothesis because they have definitive terminal values [57][12][9]. Under weak-form efficiency, no strategy based on historical prices should yield positive expected returns after the vig [69][56]. Most studies find odds are generally accurate and hard to beat, yet **many academic papers reject strict weak-form efficiency** [69][9][11].

The sources disagree on the degree and persistence of inefficiency:

- Some conclude markets are efficient within transaction costs — biases exist but cannot be exploited after the margin [26].
- Others identify persistent, profitable inefficiencies in specific European leagues [70][71].
- A significant cautionary strand argues that many reported "inefficiencies" are **statistical noise**: simulations show that statistically significant anomalies appear even in fully efficient markets

due to short samples, and over 14 seasons across major leagues, single-season inefficiencies are not persistent or systematic [18].

The research also notes that European football betting turnover exceeds €40 billion annually, that placing a bet is conceptually equivalent to purchasing a financial asset with a fixed deadline, and that peer-to-peer exchanges like Betfair mirror financial market microstructure with decentralized price formation, continuous order matching, statistical arbitrage, and cash-out features that function like stop-loss orders [72][73][74].

The Closing Line as Truth-Teller

A recurring theme: the **closing line is the market's most accurate consensus estimate** because it absorbs news, lineup changes, and betting pressure until kickoff [10][11]. This makes **Closing Line Value (CLV)** — consistently securing odds better than the closing price — the field's preferred benchmark for skill, distinguishing process from luck [9][11][28]. Because variance dominates any single result, academic consensus holds that the frequency and magnitude of beating the closing line is a more stable indicator of long-term process than raw profit and loss [11].

Efficiency is also framed as a **timing problem**: the market is often unbeatable at the closing bell but more vulnerable earlier in the day or during high-volatility periods such as injury news or lineup announcements [11]. Forecast quality does not always improve monotonically as game time approaches — some research found weekend day-game forecasts less accurate than those 90 minutes earlier, indicating scheduling-related processing inefficiencies [9]. Line movements also show **significant negative autocorrelation**, suggesting markets occasionally overreact and then correct, creating exploitable windows [9][10].

Line Movement, Sharp Money, and Syndicates

Sportsbook risk management is a dynamic, data-driven process that balances theoretical margins against market volatility [62]. Retail (soft) books often "draft" off market-making books like Pinnacle or Circa — when a sharp book moves, retail books follow to avoid being "picked off" by arbitrageurs [75][76]. **Betting syndicates** operate like investment firms, using data scientists, proprietary algorithms, and large bankrolls to exploit soft lines [77][78].

Key signals and mechanics:

- **Steam moves**: rapid, synchronized line shifts across multiple books, signaling a high-confidence syndicate position [79][80][81][82].
- **Reverse line movement (RLM)**: the line moves against the majority of public wagers, widely considered one of the strongest indicators of sharp money [79][81][80][83]. It is detected by comparing the percentage of **tickets** (number of bets) against the percentage of **handle** (dollar volume): low ticket share but high dollar share signals sharp backing [84][85].
- **Key numbers**: in sports like the NFL, lines move through key numbers (3, 7, 10); a shift from -3 to -3.5 is more significant than other moves because it crosses a high-frequency final-score differential [79][81][86]. (This is a US-sports illustration in the sources; its direct soccer analogue is less emphasized.)

Operators counter sharp action with **dynamic betting limits, cross-platform account linking to detect runners/"beards," market segmentation by risk level, and human trader overrides during high-volatility events** [77][62]. Critically, the industry compensates for identified inefficiencies through **discriminatory practices against successful clients, such as limiting accounts** of bettors who consistently exploit weaknesses [87]. **So what:** Even a genuine edge faces an execution ceiling — winning bettors get limited, so scalability and account longevity are real constraints on any strategy [87].

The research warns against naively chasing movement: once a line has moved the value may already be extracted, sharps sometimes employ "dummy moves" to mislead the market, and not all movement is sharp-driven (injuries, weather, personnel matter) [88][84][89][90][81][82].

Where the Edges Actually Are

Synthesizing the efficiency findings, exploitable edges cluster in identifiable places:

- **Lower-tier and less liquid leagues**, where pricing is softer than top-tier markets priced by sharp books [13][27].
- **Best-odds selection across multiple bookmakers** — average market odds are often efficient, but selecting the best available price can uncover inefficiency, though usually insufficient to fully overcome the vig [18][12][91].
- **Mid-table or promoted teams** where public perception skews odds [27].
- **Specific information not yet incorporated**, such as social media signals or in-play developments like goals and red cards not instantly reflected in live prices [12].
- **In-play markets** calibrated to pre-match odds — one study reported approximately **4.5% ROI** betting against Betfair in-play odds using market-calibrated Weibull AFT models [22][92].

Arbitrage, by contrast, is essentially absent: opportunities where summed inverse odds fall below 1 appear in **less than 0.1% to 0.5% of matches** [12]. Realistic value strategies therefore depend on statistical mispricing, not riskless arbitrage [12].

Market characteristic	Efficiency level	Edge potential
Top-tier 1X2 (Pinnacle-priced)	High / sharp	Low [13]
Lower-tier leagues	Lower	Higher [13][27]
Closing line (any market)	Highest	Minimal [10][11]
Early lines / pre-lineup	Lower	Higher (timing) [11]
In-play, market-calibrated	Variable	~4.5% ROI reported [22][92]
Pure arbitrage	Near-perfect	<0.5% of matches [12]

> **Bottom line for this theme:** Treat the closing line as the price you must beat and CLV as your scoreboard. The realistic edge is thin, intermittent, and concentrated in less liquid markets

> and early/volatile windows — and even then, account limiting caps how much you can deploy [11][87].

Building a Custom Model: Calibration, Devigging, and Disciplined Execution

The research converges on a clear, ordered playbook for anyone building a model to compete against bookmaker lines. The recurring message: the algorithm is the easy part; the discipline is the edge.

Step 1 — Data and Leakage Avoidance

Data leakage is identified as the dominant cause of model failure [17][29]. Leakage occurs whenever a feature contains information unavailable at prediction time:

- Using **half-time result to predict full-time result** is leakage because the two are inherently linked [17].
- Including a **game's own result in a team's cumulative season record** is leakage; it must be corrected by subtracting the current game's outcome from the running total [29].

The defense is a strict **walk-forward approach**: train only on historical data available prior to the match being predicted [27][28][29]. Temporally separating every feature from the target is non-negotiable [17]. High-quality data sources matter — Opta or FBref event data combined with minute-by-minute market data is described as essential for calibrating in-play models [22].

Step 2 — Feature Engineering Grounded in Theory

Effective features include rolling form, xG trends, venue-split performance, shots on target, corners, and attack strength [15][55][32][17]. The discipline is to **ground features in sports theory rather than p-hacking**: avoid testing arbitrary variables (like days of the week) until one appears significant by chance; anchor filters in concepts like Expected Goals [28]. As established, feature quality typically outweighs algorithm complexity [15][55][27].

Step 3 — Incorporate and Devig Market Odds

Because bookmaker odds are the strongest single predictor and ignoring them leaves you at a KL-disadvantage, market prices should enter the model directly [21][22][23]. One documented approach: invert 1X2 bookmaker odds (using the Skellam distribution) to recover implied home and away Poisson means, then combine them with historical attack-defense estimates via **Bayesian weighting** [22].

Recovering "true" probabilities from quoted odds — **devigging** — is itself a modeling decision with several documented methods [16]:

- **Naive normalization**: divide each implied probability by the sum of all implied probabilities (1 + margin) to force them to sum to 1 [60][16].

- **Additive and multiplicative models**: alternative frameworks to strip the margin [93].
- **Shin method**: accounts for informed players and risk aversion [16].
- **Knowles-Woodland** and **Rascher's method** (closed-form for two-outcome markets) [16].
- **Domain-driven identification**: integrating multiple betting markets simultaneously (e.g., via Dixon-Coles) to leverage "smart liquidity" across outcome groups, providing more robust calibration [16].

Because margins are not uniform (the FLB skew), naive normalization can misstate true probabilities — method choice should be tested, with log-loss and Brier score as comparative metrics [24][16][25].

Step 4 — Calibrate the Probabilities

This is the step the research elevates above all others. **Calibration is more important than accuracy for profitability** [5][6][8]. A miscalibrated, overconfident model inflates the perceived edge, leading to excessive Kelly stakes and elevated risk of ruin [7]. Recalibration techniques map raw outputs to true frequencies:

- **Platt scaling**: fits a sigmoid/logistic function to model outputs; fast but may underfit complex distributions and is effective for SVMs [7][94][95].
- **Isotonic regression**: non-parametric monotonic mapping; powerful with large datasets but prone to overfitting on small samples, and may not improve Random Forests without ample calibration data [15][7][96][95].
- **Beta calibration**: parametric, effective for tail-heavy / extreme probability skews [7][94][97].
- **Temperature scaling**: divides logits by a learned scalar, common in neural networks [7].
- **Venn-Abers predictors**: distribution-free with validity guarantees under exchangeability; strong at reducing log-loss [97][95].

The research cautions there is **no one-size-fits-all method** — for high-performing modern tabular models, Platt scaling and isotonic regression can occasionally degrade proper scoring performance, and the calibration-refinement decomposition shows optimizing one can hinder the other [97][98]. Practical advice: select the calibration method with the lowest Brier score on a held-out calibration set [95].

Step 5 — Evaluate with Proper Scoring Rules

The research is unanimous that classification accuracy is the wrong metric. The correct toolkit:

- **Ranked Probability Score (RPS)**: the standard benchmark for win/draw/loss forecasts, computed as RPS = (1/(r-1)) * sum over i of (cumulative(p_j - a_j))^2, where r is the number of outcomes, p_j the predicted probability, and a_j the actual outcome; lower is better [99][100]. Its virtue is **sensitivity to ordinal distance** — a draw is "closer" to a home win than an away win, so misplaced probability on a distant outcome is penalized more heavily [101][102]. Some researchers argue this distance-sensitivity does not improve practical utility and that local scoring rules like the **ignorance score** may evaluate better [101].

- **Log loss (cross-entropy)**: -(1/N) * sum of [ylog(p) + (1-y)log(1-p)]; ranges 0 to infinity, penalizes confident-but-wrong predictions severely, and is favored as a training objective — but assumes reasonably calibrated probabilities [103][7].
- **Brier score**: mean squared difference between predicted probability and outcome, (1/N)*sum of (y-p)^2; ranges 0 to 1, robust to outliers, captures both calibration and discrimination [103][7][94][98].
- **Expected Calibration Error (ECE)** and **reliability diagrams**: diagnostic tools plotting predicted probability against observed frequency across bins [7][94].
- A novel evaluation treats models as **canonical Kelly bettors** in an iterative contest where bankroll proxies Bayesian credibility, allowing dynamic model comparison before outcomes are known — reportedly more accurate than static log-loss or Brier at distinguishing correct models, especially when models exhibit faulty recency bias or non-predictive variables [104][105].

Step 6 — Backtest Without Illusions

Backtesting is prone to the "backtesting illusion" if it ignores market dynamics [28]. Requirements drawn from the research:

- **Out-of-sample testing**: split into training and test sets; consistent performance across both indicates robustness, while a test-set collapse signals overfitting [28][29].
- **Adequate sample size**: typically **250–500+ bets** to overcome football's high variance [28].
- **Benchmark to the closing price**: identify CLV as a reliable indicator of long-term profitability rather than assuming you can get arbitrary historical odds [28].
- **Assess maximum drawdown**: a paper-profitable strategy is useless if the bankroll cannot survive a prolonged losing streak [28].

One important model-selection caveat: **Reinforcement Learning is often unsuitable for match prediction** because, unlike chess or financial markets, soccer matches are largely independent events where current outcomes do not influence future states [27].

Step 7 — Value Identification and Bankroll Discipline

Profitability comes from the **edge** — the gap between the model's estimated probability and the market-implied probability [106][27]. The standard staking framework is the **Kelly criterion**, which maximizes the long-term logarithmic growth rate of the bankroll [33][107][34]. In simple form, f = p - (1-p)/o, where p is the estimated win probability, o the decimal odds, and f the optimal fraction to wager [104]. Edge can be expressed as (odds × p) - 1 [7].

But full Kelly is dangerous in practice:

- It assumes the bettor knows the **true** probability, which is rarely possible in sports [33].
- Full Kelly produces extreme volatility and risk of "gambler's ruin" during losing streaks [34][35].
- Real markets impose payout caps and limited liquidity that make aggressive Kelly suboptimal [34].

Hence the field standard is **fractional Kelly** — half or quarter — to reduce volatility and cushion probability-estimation error [15][27][33]. More advanced variants include **distributionally robust Kelly**, which optimizes worst-case growth across a set of possible probability distributions, and **adaptive strategies** that adjust the fraction with model confidence [108][33]. Multivariate Kelly optimization manages portfolios of multiple simultaneous, possibly correlated bets [109].

A final sobering reality: naive strategies such as **blindly betting favorites almost universally lose money** because of the margin, and profitable models tend to find success precisely where liquidity is thin or public perception skews odds [27].

Build stage	Key tool / metric	Primary risk if skipped
Data / leakage	Walk-forward, feature audit	Inflated, unreproducible backtests [17][29]
Feature engineering	xG, form, ratings; theory-grounded	p-hacking, overfit [28][15]
Market integration	Bayesian odds combination	KL-disadvantage vs market [21][22]
Devig	Shin / multi-market / normalization	Misstated true probabilities [16][25]
Calibration	Platt, isotonic, Beta, Venn-Abers	Inflated edge, over-staking [7][5]
Evaluation	RPS, log-loss, Brier, ECE	Optimizing wrong objective [99][103]
Backtest	Out-of-sample, 250–500+ bets, CLV, drawdown	Backtesting illusion [28]
Staking	Fractional Kelly	Bankroll ruin [34][33]

> **Bottom line for this theme:** A profitable pipeline is a disciplined sequence — leak-free data, theory-grounded features, market-integrated and devigged probabilities, rigorous calibration, proper-scoring evaluation, honest backtesting against the closing line, and fractional-Kelly staking. Skip any one step and the edge evaporates. Calibration and CLV are the two non-negotiables [5][28][7].

Recommendations

1. **Quants and model builders: select and tune every model on calibration metrics, not accuracy.** Use Brier score, log-loss, RPS, ECE, and reliability diagrams as your selection criteria. Trigger condition: if your model selection process currently optimizes classification accuracy or "winner picked" rates, change it before placing a single bet — the evidence shows the difference between +34.69% and -35.17% ROI hinges on this choice [5][6][8].
2. **Anyone integrating market data: combine bookmaker odds into your model via Bayesian weighting rather than competing against them from scratch.** Invert 1X2 odds (e.g., through the Skellam distribution) to recover implied scoring means and blend with your attack-defense

estimates [22]. Trigger: adopt this whenever your model's standalone RPS fails to beat the devigged market — which the research suggests will be most of the time [21][23].

3. **Edge-hunters: concentrate effort in less liquid, lower-tier leagues and early/volatile pricing windows, and shop for best odds across books.** Avoid the top-tier 1X2 markets priced by sharp books and avoid parlays where margin compounds [13][27][58]. Trigger: prioritize a market when your devigged model probability diverges materially from the best available price and the market is demonstrably softer than top-tier benchmarks [12][91].

4. **All bettors: instrument CLV tracking as your primary KPI from day one.** Record the closing line for every wager and measure how often and by how much you beat it [28]. Trigger: if your CLV is not consistently positive over a 250–500+ bet sample, your process has no demonstrated edge regardless of short-run profit [9][11][28].

5. **Engineers: enforce walk-forward validation and audit every cumulative feature for self-inclusion.** Remove half-time-to-full-time leakage and subtract current-game results from running totals [17][29]. Trigger: re-audit any feature whose addition produces a suspiciously large backtest improvement — it is more likely leakage than signal [17].

6. **Risk managers / bettors: stake fractional Kelly (quarter to half) and stress-test maximum drawdown.** Never deploy full Kelly with model-estimated probabilities [15][33][34]. Trigger: reduce the Kelly fraction further whenever calibration confidence is low or recent reliability diagrams show drift [35][108].

7. **Strategists: plan for account limiting as an operational constraint, not an afterthought.** Successful clients are limited by operators [87]. Trigger: as soon as a strategy demonstrates positive CLV at scale, build account diversification and stake-distribution plans because pricing power and access will shift against you [77][87].

8. **Practitioners building from the statistical track: use Dixon-Coles as the mandatory baseline benchmark.** Any added complexity (bivariate Poisson, hierarchical Bayesian, ML ensembles) must demonstrably lower RPS against this baseline to justify itself [1][4]. Trigger: discard model complexity that does not beat Dixon-Coles out-of-sample [4].

Caveats & Limitations

Conflicting evidence on market efficiency. The sources genuinely disagree on whether soccer markets are exploitable. Some report persistent, profitable inefficiencies in specific European leagues [70][71]; others conclude inefficiencies vanish after transaction costs [26]; and a strong cautionary strand argues many reported anomalies are statistical noise that disappears over large samples (14 seasons across major leagues) [18]. This report presents all positions rather than adjudicating among them. Readers should treat any claimed edge skeptically until validated over large out-of-sample data.

Profitability claims are limited and context-specific. The most concrete positive result — approximately 4.5% ROI — comes from market-calibrated Weibull AFT models bet against Betfair in-play odds in particular studies [22][92]. This is not a general guarantee; the "KellyBench" research shows even frontier AI models often fail to profit over a full season due to non-stationarity and the overround [19]. Single-study ROI figures should not be extrapolated to other leagues, periods, or market conditions.

Some pricing illustrations are drawn from US sports. Concepts like "key numbers" (3, 7, 10) and the significance of a -3 to -3.5 shift are explicitly NFL/point-spread phenomena in the sources [79][81][86]; their direct applicability to soccer 1X2 and Asian handicap markets is not established in the research and should be treated as analogy, not fact.

Calibration methods are not universally beneficial. The research explicitly warns that Platt scaling and isotonic regression can degrade proper-scoring performance for some high-performing modern models, that isotonic regression overfits small samples, and that there is no one-size-fits-all method [96][97][98]. Method choice must be validated empirically on a held-out set [95].

Devig method ambiguity. Multiple devig methods (naive normalization, Shin, Knowles-Woodland, Rascher, domain-driven) coexist, and the sources note they yield different "true" probabilities because margins are non-uniform [16][25]. There is no single agreed-upon correct method, so measured "edge" is partly a function of the devig choice — a material modeling uncertainty.

Data and source granularity. Much xG and tracking research relies on open-source datasets that may limit generalizability, while proprietary data held by professional clubs is unavailable to most researchers [53][4]. Several search results reference academic profiles or sources that did not contain substantive content on the queried topic (notably one scholar.google.com search that returned unrelated profiles) [110][111][112]; no claims in this report rest on those.

Forward-looking and structural constraints. Claims that bookmakers "continuously improve forecast precision" and that markets are getting tighter are general trends in the literature, not precisely quantified facts [18][113]. Finally, the structural reality of account limiting means that even a statistically validated edge faces a practical ceiling on deployable capital that the academic ROI figures do not capture [87].

> **Final word for decision-makers:** The methods in this report are real, open, and well-documented — but so are the obstacles. The defensible strategy is not a secret model; it is superior discipline in calibration, devigging, leakage control, CLV measurement, and risk sizing, executed against the humbling assumption that the closing line already knows almost everything you do [11][5][28].

References

1. https://grokipedia.com/page/DixonColes_model
2. https://onlinelibrary.wiley.com/doi/abs/10.1111/j.1467-9574.1982.tb00782.x
3. http://www2.stat-athens.aueb.gr/~jbn/papers2/08_Karlis_Ntzoufras_2003_RSSD.pdf
4. https://pub.towardsai.net/the-algorithm-that-ruined-my-ucl-final-30dbbe110c09
5. https://www.sciencedirect.com/science/article/pii/S266682702400015X

6. https://arxiv.org/abs/2303.06021
7. https://www.sports-ai.dev/blog/ai-model-calibration-brier-score
8. https://papers.ssrn.com/sol3/Delivery.cfm/9d490e32-e995-47d6-b18d-ba041050d2e1-MECA.pdf?abstractid=4705918&mirid=1
9. https://www.researchgate.net/publication/372441761_Inefficient_Forecasts_at_the_Sportsbook_An_Analysis_of_Real-Time_Betting_Line_Movement
10. https://w4.stern.nyu.edu/finance/docs/pdfs/Seminars/moskowitz_seminar_paper.pdf
11. https://joesaumarez.co.uk/sports-betting-market-efficiency-and-the-closing-line
12. https://www.researchgate.net/publication/328072253_Efficiency_of_online_football_betting_markets
13. https://www.researchgate.net/publication/351985837_Favorite-Longshot_Bias_and_Market_Efficiency_in_the_Soccer_Betting_Market
14. https://www.sciencedirect.com/science/article/pii/S0169207024000670
15. https://journals.sagepub.com/doi/10.1177/22150218261416681
16. https://www.mdpi.com/2227-7390/13/24/3976
17. https://medium.com/@olufaithy/predictive-modelling-for-football-match-results-a-python-and-machine-learning-approach-ec94e9fa2fef
18. https://journals.sagepub.com/doi/10.1177/15270025231204997
19. https://arxiv.org/html/2604.27865
20. https://dashee87.github.io/football/python/predicting-football-results-with-statistical-modelling-dixon-coles-and-time-weighting/
21. https://arxiv.org/pdf/2107.08827
22. https://arxiv.org/html/2605.16066
23. https://arxiv.org/pdf/2307.10411
24. https://www.researchgate.net/publication/368757288_Calculating_The_Bookmaker
25. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4374924
26. https://www.sciencedirect.com/science/article/abs/pii/S1062976916000041
27. https://medium.com/@omnimahui/profitable-england-premier-league-betting-strategy-a-full-end-to-end-experiment-2a53b32ba16d
28. https://www.predictology.co/blog/are-you-making-these-common-backtesting-mistakes-how-to-validate-your-football-system/
29. https://dhruviyer.substack.com/p/predicting-nfl-outcomes-using-deep
30. https://www.researchgate.net/figure/Scatter-Plot-of-the-Dixon-Coles-Strength-Parameters-for-English-Premier-League-Teams_fig4_338028907
31. https://arxiv.org/html/2505.01902v1
32. https://link.springer.com/article/10.1186/s40537-024-01008-2
33. https://ar5iv.labs.arxiv.org/html/2107.08827
34. https://ar5iv.labs.arxiv.org/html/1701.01427
35. https://arxiv.org/pdf/2303.06021
36. https://www.academia.edu/118422937/Application_of_Poisson_and_Dixon_Coles_models_on_football_match_outcome_prediction_and_research_of_a_positive_return_over_investment_in_betting_market
37. http://www.90minut.pl/misc/maher.pdf
38. https://www.slideshare.net/slideshow/1982-maher-modelling-association-football-scores/24914293
39. https://www.semanticscholar.org/paper/Modelling-association-football-scores-Maher/1bed7ee4bfe34d64b7a96e640ee5b9e7cb3d1376
40. https://www.scribd.com/document/414405327/maher-pdf
41. https://pdfs.semanticscholar.org/33e1/ca44b495df0b419be79daac85d2527a3bc8f.pdf
42. https://opisthokonta.net/?cat=48
43. https://docs.pena.lt/y/models/dixon_coles.html
44. https://arxiv.org/pdf/2307.02139
45. https://www.researchgate.net/publication/227719079_Analysis_of_sports_data_using_bivariate_Poisson_models
46. http://www2.stat-athens.aueb.gr/~karlis/Bivariate%20Poisson%20Regression.pdf
47. https://www.researchgate.net/publication/228621612_Bayesian_modelling_of_football_outcomes_Using_the_Skellam
48. https://www.academia.edu/1008427/Bayesian_modelling_of_football_outcomes_using_the_Skellams_distribution_for_the_goal_difference
49. http://eprints.utar.edu.my/4981/1/WONG_ZHEN_PING.pdf

50. https://www.academia.edu/119174296/Predicting_football_scores_via_Poisson_regression_model_applications_to_the_National_Football_League
51. https://arxiv.org/pdf/1710.00001
52. https://www.researchgate.net/publication/362022551_Dissecting_Poisson_based_prediction_models_in_association_football_A_comprehensive_look_at_methodology_assumptions_and_accuracy_using_data_from_the_main_European_Leagues_2011_-_2022
53. https://www.ajbuckeconbikesail.net/wkpapers/Airports/MVPoisson/soccer_betting.pdf
54. https://arxiv.org/html/2410.21484v1
55. https://www.sciencedirect.com/science/article/pii/S2772662224001413
56. https://myweb.ecu.edu/robbinst/PDFs/Weak%20Form%20Efficiency%20in%20Sports%20Betting%20Markets.pdf
57. https://www.reading.ac.uk/web/files/economics/emdp201910.pdf
58. https://en.wikipedia.org/wiki/Mathematics_of_bookmaking
59. https://www.gammastack.com/blog/how-do-bookmakers-generate-sports-odds/
60. https://econ.berkeley.edu/sites/default/files/siwei-xu_jasmine.pdf
61. https://www.ucd.ie/economics/t4media/WP23_04.pdf
62. https://altenar.com/en-us/blog/house-edge-in-sports-betting-how-the-best-sportsbooks-manage-their-edge/
63. https://www.erfcinc.org/group-page/educational-resource-group/discussion/e1b58736-7fb7-4163-8f37-5e597856ce57
64. https://onlinelibrary.wiley.com/doi/10.1111/ecca.12500
65. https://academic.oup.com/oep/article/78/1/90/8244336
66. https://onlinelibrary.wiley.com/doi/10.1111/sjpe.70017
67. https://www.researchgate.net/publication/368848754_Forecasting_Soccer_Matches_With_Betting_Odds_A_Tale_of_Two_Markets
68. https://www.ucd.ie/economics/t4media/WP23_05.pdf
69. https://www.researchgate.net/publication/365618651_On_the_Efficiency_of_Sports_Betting_Markets
70. https://www.researchgate.net/publication/4920285_The_Favourite-Longshot_Bias_and_Market_Efficiency_in_UK_Football_Betting
71. https://www.sseriga.edu/sites/default/files/2024-05/8Paper_Gabor_Paju.pdf
72. https://arxiv.org/html/2505.21275v1
73. https://arxiv.org/html/2605.30209v1
74. https://arxiv.org/html/2409.13528
75. https://help.outlier.bet/en/articles/9922960-how-sportsbooks-set-odds-soft-vs-sharp-books
76. https://edgeslip.com/articles/sportsbook-screen
77. https://oddsmatrix.com/betting-syndicate/
78. https://www.europeanbusinessreview.com/how-bet-syndicates-influence-odds-in-major-sportsbooks/
79. https://oddsindex.com/guides/line-movement-betting
80. https://theedgesniper.com/blog/sportsbook-trends
81. https://www.deucescracked.com/sports-betting/strategy/line-movement
82. https://www.predictem.com/betting/strategy/line-movement-analysis/
83. https://www.sportingnews.com/us/betting/news/how-sportsbooks-set-betting-lines/06bbb4dfdd9b14f04abed737
84. https://www.fantasylife.com/articles/betting/what-is-reverse-line-movement
85. https://www.sportbotai.com/blog/sharp-money-explained-sports-betting
86. https://www.boydsbets.com/sharp-line-moves/
87. https://www.researchgate.net/publication/276014117_Are_Sports_Betting_Markets_Prediction_Markets_Evidence_From_a_New_Test
88. https://www.betsperts.com/news/sports-betting-guides/sharp-money-trackers/
89. https://www.actionnetwork.com/education/reverse-line-movement
90. https://www.betsperts.com/news/sports-betting-guides/reverse-line-movements/
91. https://www.researchgate.net/publication/291191792_Efficiency_in_Betting_Markets_Evidence_From_English_Football
92. https://arxiv.org/pdf/2605.16066
93. https://outlier.bet/wp-content/uploads/2023/08/2017-clarke-adjusting_bookmakers_odds.pdf
94. https://arxiv.org/html/2509.23665v1
95. https://arxiv.org/pdf/2409.04874
96. https://arxiv.org/html/2501.16756v1
97. https://arxiv.org/html/2601.19944
98. https://arxiv.org/html/2501.19195v1
99. https://arxiv.org/html/2309.14807

100. https://arxiv.org/pdf/1908.08980
101. https://arxiv.org/abs/1908.08980
102. https://arxiv.org/html/2508.05891v1
103. https://www.dratings.com/log-loss-vs-brier-score/
104. https://arxiv.org/html/2602.09982v1
105. https://arxiv.org/pdf/2602.09982
106. https://www.footixify.com/
107. https://arxiv.org/pdf/2010.12508
108. https://arxiv.org/pdf/1812.10371
109. https://arxiv.org/pdf/2604.24723
110. https://scholar.google.com/citations?hl=en&user=OxZYSkoAAAAJ
111. https://scholar.google.com/citations?hl=en&user=OsjprIsAAAAJ
112. https://scholar.google.com/intl/en/scholar/inclusion.html
113. https://research-api.cbs.dk/ws/portalfiles/portal/60750333/237159_final_digital.pdf